

Task 1.1

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms

class CNN(nn.Module):

    def __init__(self):
        super(CNN, self).__init__()

        self.layer1 = nn.Sequential(
            nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))

        self.layer2 = nn.Sequential(
            nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0),
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2))

        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        out = self.layer1(x)
        out = self.layer2(out)

        out = out.view(-1, 16 * 5 * 5)

        out = F.relu(self.fc1(out))
        out = F.relu(self.fc2(out))
        out = self.fc3(out)

        return out
```

Task 1.2

Output Size:

(64 x 64 x 1): Activation size = $64 * 64 * 1 = 4096$

Layer_1: $64 - 5 + 1 = 60 \times 60 \times 6$

MaxPool: $(60 - 2) / 2 + 1 = 30 \times 30 \times 6$

Layer_2: $30 - 5 + 1 = 26 \times 26 \times 16$

MaxPool: $(30 - 2) / 2 + 1 = 15 \times 15 \times 16$

fc1: 120×1

fc2: 84×1

fc3: 20×1

Number of Parameters:

Layer1:

Weights = $5^2 \times 1 \times 6 = 150$

Biases = 6

Parameters = $150 + 6 = 156$

Layer2:

Weights = $5^2 \times 1 \times 16 = 400$

Biases = 16

Parameters = $400 + 16 = 416$

fc1 connected to a Conv Layer:

Weights: $15^2 \times 16 \times 120 = 432000$

Biases: 120

Parameters = 432120

fc2 connected to the previous fc layer:

Weights: $120 \times 84 = 10080$

Biases: 84

Parameters: 10164

fc3 connected to the previous fc layer:

Weights: $84 \times 10 = 840$

Biases: 10

Parameters: 850

Total Parameters: $156 + 416 + 432120 + 10164 + 850 = 443706$

Task 1.3

```
def train(self):
    trainloader, _ = self.load_data()
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.SGD(cnn.parameters(), lr=0.001, momentum=0.9)

    for epoch in range(10):
        for i, data in enumerate(trainloader, 0):
            inputs, labels = data

            optimizer.zero_grad()
            outputs = cnn(inputs)

            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

def eval(self):
    _, testloader = self.load_data()
    total_correct = 0
    total_images = 0

    with torch.no_grad():
        for data in testloader:
            images, labels = data
            outputs = cnn(images)
            _, predicted = torch.max(outputs.data, 1)

            total_images += labels.size(0)
            total_correct += (predicted == labels).sum().item()

    model_accuracy = total_correct / total_images * 100
    print('Model accuracy on {0} test images: {1:.2f}%'.format(total_images,
model_accuracy))
```

```

def load_data(self):
    transform = transforms.Compose(
        [
            transforms.ToTensor()
        ])

    trainset = torchvision.datasets.CIFAR10(root='./Data',
                                             train=True,
                                             download=True,
                                             transform=transform)

    trainloader = torch.utils.data.DataLoader(trainset,
                                              batch_size=16,
                                              shuffle=True)

    testset = torchvision.datasets.CIFAR10(root='./Data',
                                             train=False,
                                             download=True,
                                             transform=transform)

    testloader = torch.utils.data.DataLoader(testset,
                                              batch_size=16,
                                              shuffle=False)

    return trainloader, testloader

if __name__ == "__main__":
    cnn = CNN()
    cnn.train()
    cnn.eval()

```

Model accuracy on 10000 test images: 58.14%

Task 1.4

```

def load_data(self):
    transform = transforms.Compose(
        [
            transforms.RandomHorizontalFlip(p=0.5),
            transforms.RandomCrop(32, padding=4),
            transforms.ToTensor(),
            transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
        ])

    trainset = torchvision.datasets.CIFAR10(root='./Data',
                                             train=True,
                                             download=True,
                                             transform=transform)

    trainloader = torch.utils.data.DataLoader(trainset,
                                              batch_size=16,
                                              shuffle=True)

    testset = torchvision.datasets.CIFAR10(root='./Data',
                                             train=False,
                                             download=True,
                                             transform=transform)

    testloader = torch.utils.data.DataLoader(testset,
                                              batch_size=16,
                                              shuffle=False)

    return trainloader, testloader

```

Input normalization:

Model accuracy on 10000 test images: 61.72%

In comparison to the baseline model its about 3% better.

Input normalization and augmentation:

Model accuracy on 10000 test images: 57.08%

In comparison to the baseline model its about 1% worse.

Task 1.5

0.1: 10.00%

0.01: 51.34%

0.001: 57.08%

0.0001: 32.60%

The learning rate tells the optimizer how far to move the weights in the direction opposite of the gradient. If it is low, then training is more reliable, but optimization will take a lot of time because steps towards the minimum of the loss function are tiny, see above: 0.1 -> 10%.

If the learning rate is high, then training may not converge. Weight changes can be so big that the optimizer overshoots the minimum and makes the loss worse, 0.0001: 32.60%. The best learning rate in the four examples is 0.001 -> 57.08%.

Task 1.6

```
def __init__(self):
    super(CNN, self).__init__()

    self.layer1 = nn.Sequential(
        nn.Conv2d(3, 6, kernel_size=5, stride=1, padding=0),
        nn.BatchNorm2d(6),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2))

    self.layer2 = nn.Sequential(
        nn.Conv2d(6, 16, kernel_size=5, stride=1, padding=0),
        nn.BatchNorm2d(16),
        nn.ReLU(),
        nn.MaxPool2d(kernel_size=2, stride=2))

    self.fc1 = nn.Linear(16 * 5 * 5, 120)
    self.fc2 = nn.Linear(120, 84)
    self.fc3 = nn.Linear(84, 10)

def forward(self, x):
    out = self.layer1(x)
    out = self.layer2(out)

    out = out.view(-1, 16 * 5 * 5)

    out = F.relu(self.fc1(out))
    out = F.relu(self.fc2(out))
    out = self.fc3(out)

    return out
```

In contrast to the baseline model with 58.14% the model with a batch normalization layer performs a accuracy of 56.33%. So it's a little bit worse, but it should reduce overfitting because it has slight regularization effects.

Task 1.7

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision.transforms as transforms
import torchvision

class Block(nn.Module):
    '''Depthwise conv + Pointwise conv'''
    def __init__(self, in_planes, out_planes, stride=1):
        super(Block, self).__init__()
        self.conv1 = nn.Conv2d(in_planes, in_planes, kernel_size=3, stride=stride,
padding=1, groups=in_planes, bias=False)
        self.bn1 = nn.BatchNorm2d(in_planes)
        self.conv2 = nn.Conv2d(in_planes, out_planes, kernel_size=1, stride=1,
padding=0, bias=False)
        self.bn2 = nn.BatchNorm2d(out_planes)

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = F.relu(self.bn2(self.conv2(out)))
        return out

class MobileNet(nn.Module):
    # (128,2) means conv planes=128, conv stride=2, by default conv stride=1
    cfg = [64, (128,2), 128, (256,2), 256, (512,2), 512, 512, 512, 512, 512,
(1024,2), 1024]

    def __init__(self, num_classes=10):
        super(MobileNet, self).__init__()
        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, stride=1, padding=1,
bias=False)
        self.bn1 = nn.BatchNorm2d(32)
        self.layers = self._make_layers(in_planes=32)
        self.linear = nn.Linear(1024, num_classes)

    def _make_layers(self, in_planes):
        layers = []
        for x in self.cfg:
            out_planes = x if isinstance(x, int) else x[0]
            stride = 1 if isinstance(x, int) else x[1]
            layers.append(Block(in_planes, out_planes, stride))
            in_planes = out_planes
        return nn.Sequential(*layers)

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.layers(out)
        out = F.avg_pool2d(out, 2)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out

    def train(self):
        trainloader, _ = self.load_data()
        criterion = nn.CrossEntropyLoss()
        optimizer = torch.optim.SGD(mobileNet.parameters(), lr=0.001, momentum=0.9)

        for epoch in range(5):
            for i, data in enumerate(trainloader, 0):
                inputs, labels = data

                optimizer.zero_grad()
                outputs = mobileNet(inputs)

                loss = criterion(outputs, labels)
```

```

        loss.backward()
        optimizer.step()

def eval(self):
    _, testloader = self.load_data()
    total_correct = 0
    total_images = 0

    with torch.no_grad():
        for data in testloader:
            images, labels = data
            outputs = mobileNet(images)
            _, predicted = torch.max(outputs.data, 1)

            total_images += labels.size(0)
            total_correct += (predicted == labels).sum().item()

    model_accuracy = total_correct / total_images * 100
    print('Model accuracy on {0} test images: {1:.2f}%'.format(total_images,
model_accuracy))

def load_data(self):
    transform = transforms.Compose(
        [
            transforms.RandomHorizontalFlip(p=0.5),
            transforms.RandomCrop(32, padding=4),
            transforms.ToTensor(),
            transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
        ])

    trainset = torchvision.datasets.CIFAR10(root='./Data',
                                           train=True,
                                           download=True,
                                           transform=transform)

    trainloader = torch.utils.data.DataLoader(trainset,
                                              batch_size=16,
                                              shuffle=True)

    testset = torchvision.datasets.CIFAR10(root='./Data',
                                           train=False,
                                           download=True,
                                           transform=transform)

    testloader = torch.utils.data.DataLoader(testset,
                                             batch_size=16,
                                             shuffle=False)

    return trainloader, testloader

if __name__ == "__main__":
    mobileNet = MobileNet()
    mobileNet.train()
    mobileNet.eval()

```

Not enough computing power to calculate and get some results...